CI/CD PW Workflow

Github Actions and PW Client

Introduction

The setup steps described here have already been done and are left in this account as a template. Please feel free to replicate these steps to get a feel for how they work.

Where is the source code?

1. weather-cluster-demo: repository with weather model install and launch scripts
2. test-workflow-action: repository containing the GitHub action linked to weather-cluster-demo.
3. beta.parallel.works: PW SaaS platform for trial. Please login here with PW credentials. The code is already on the platform in a workflow directory in `/pw/workflows/weather-cluster-demo`.

TODO:

- Add connectivity driver information (no need to setup EFA/gVINC yourself)
- Add hpc6a and build/test everything on Azure

Introduction

Where is the documentation?

1. weather-cluster-demo/README.md: Software installation and how to run the weather model application.
2. test-workflow-action/README.md: using the GitHub action
3. This slide deck: Summary of the steps to setup the model launched by action.
4. PW platform  buttons will open a new tab (but slightly out of date).

Where to start?

1. Log in to PW to view the resource configurations, IDE, starting/stopping clusters, and interactive access to `*.clusters.pw`.
2. The GitHub action can be run directly from weather-cluster-demo on GitHub - the same weather model will be run on atNorth, AWS, and GCE.
3. **On PW,** `/pw/workflows/weather-cluster-demo/main.sh` is the core code launched by the workflow; it will clone the repo, launch the model on the clusters, and monitor the status of the application.

Introduction

Where are things on the clusters?

1. The WRF application code is in `/var/lib/pworks/spack` on GCE and AWS. atNorth, the application code is in `/shared/wrf/spack`.
2. GCE and AWS clusters share `$HOME` between the head node and worker nodes, so the working directory for WRF is in `$HOME/weather-forecast-demo/<jobid>/weather-forecast-demo/conus_12km` . Initial setup of this working directory is done with `local_setup.sh` because `$HOME` is not persistent (i.e. not in a cloud disk or in the image).
3. atNorth clusters do not share `$HOME` between head node and worker nodes and the `$HOME`(s) are persistent. Instead, the working directory is in `/shared/weather-forecast-demo/<jobid>/weather-forecast-demo/conus_12km` .
4. In all cases, you can track the WRF run via the main log file `conus<job>*.out`, and the 0 rank MPI process' log `rsl.error.0000`.
5. Output is in NetCDF format in `wrfout*` files.

Setup step 1: Setup projects

Projects do **NOT** apply to atNorth

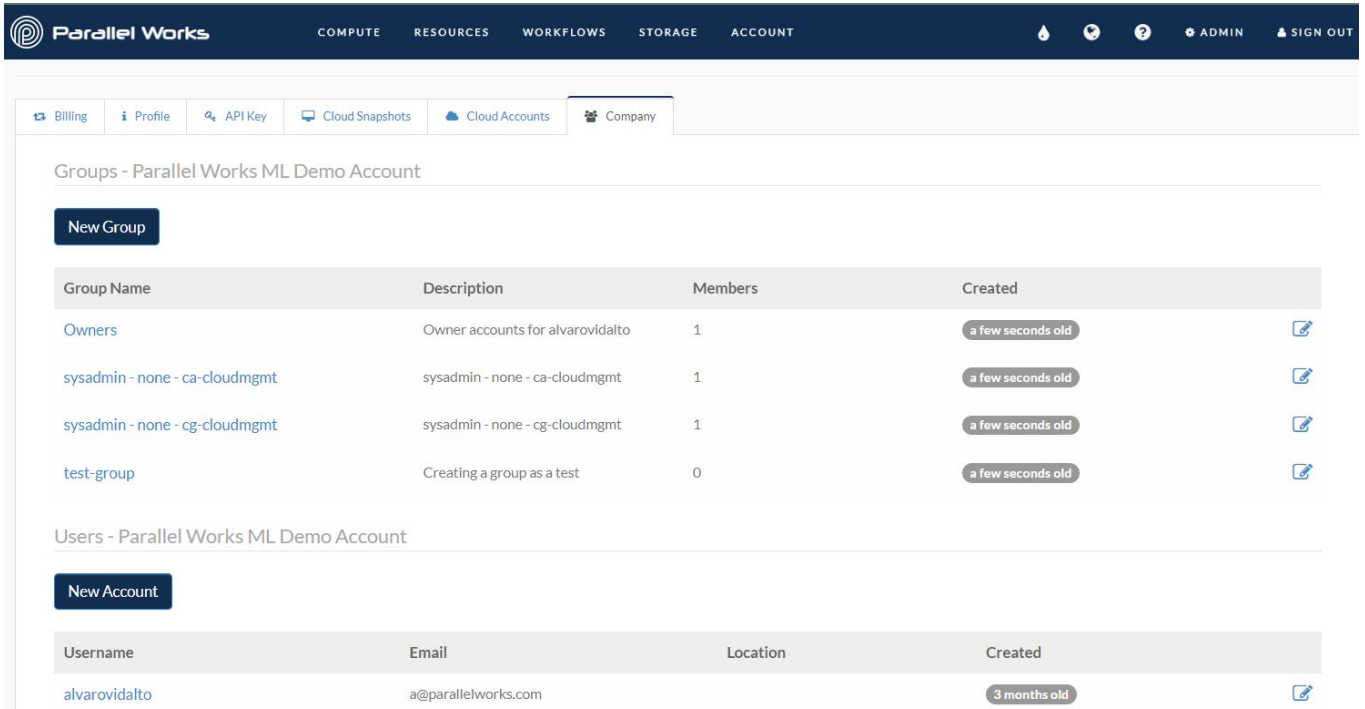

PW “main” user accounts can create subaccounts and groups of accounts in the Account -> Company tab. These groups are the “projects” used in the cluster configuration step, later. To use a group/project, simply add a user to it. This applies to main user accounts and subaccounts. Currently, projects for the major cloud providers require the following prefixes:

- GCE: cg-<project_name>
- AWS: ca-<project_name>
- Azure: cz-<project_name>

Setup step 2a: Manage images

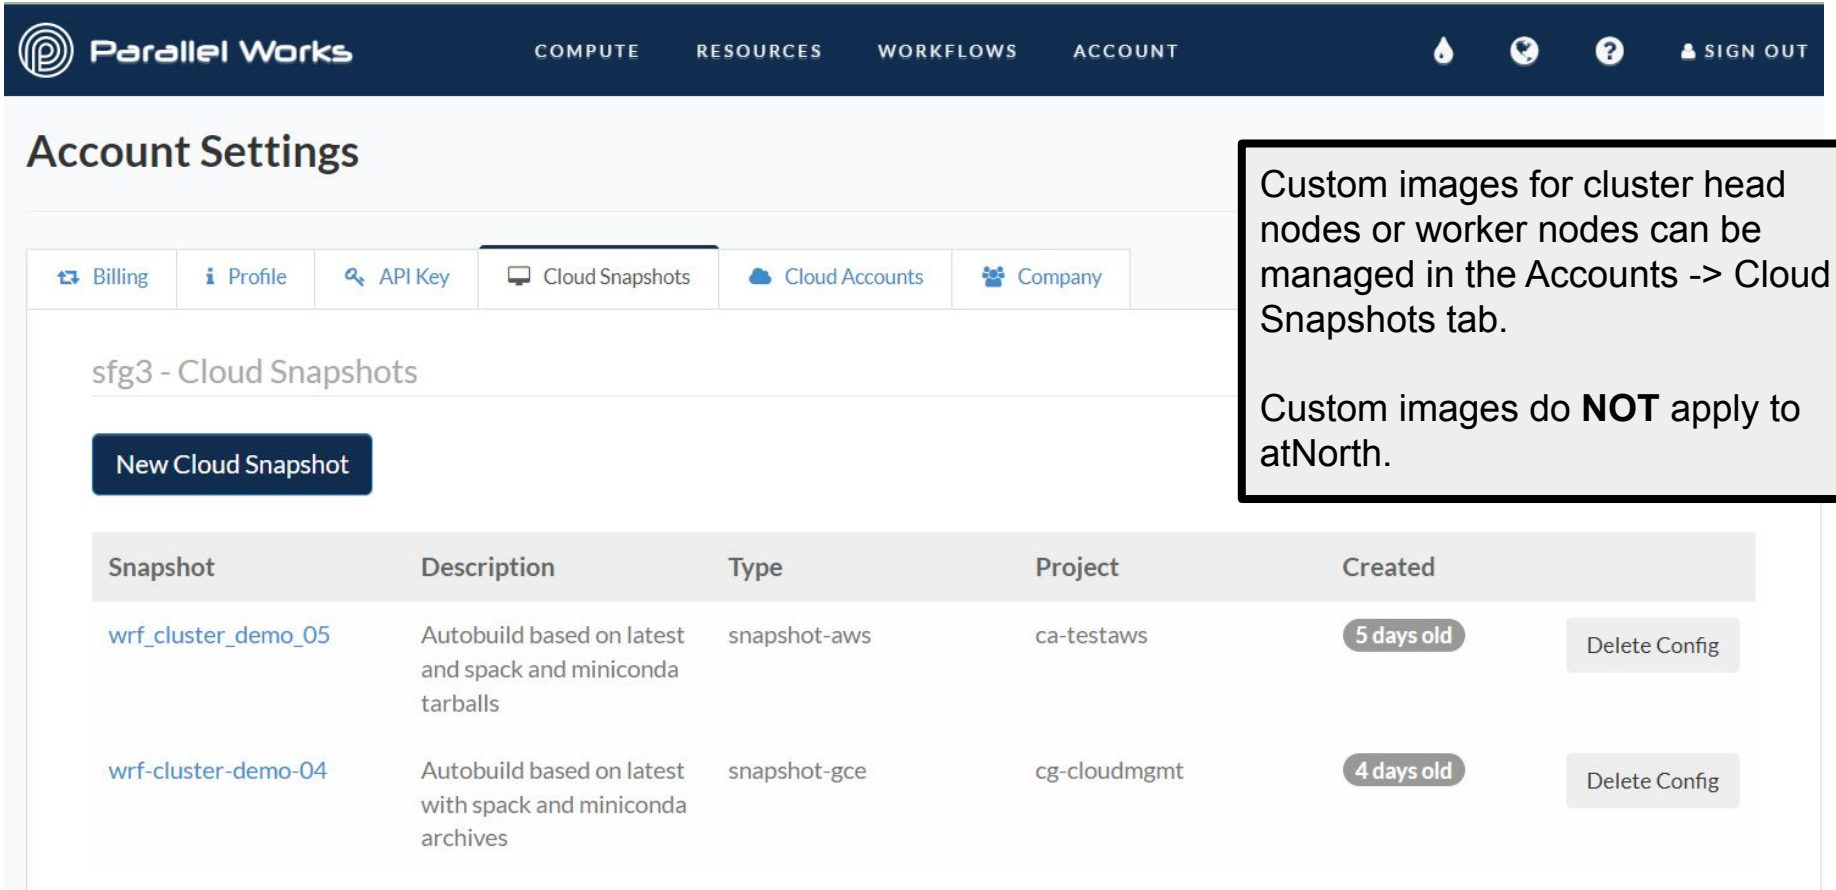

Custom images for cluster head nodes or worker nodes can be managed in the Accounts -> Cloud Snapshots tab.

Custom images do **NOT** apply to atNorth.

Setup step 2b: Build images

Parallel Works | COMPUTE | RESOURCES | WORKFLOWS | ACCOUNT | SIGN OUT

Account Settings

Billing | Profile | API Key | Cloud Snapshots | Cloud Accounts | Company

WRF_CLUSTER_DEMO_05 Snapshot Settings

Type: Amazon Web Services

Project: ca-testaws

Base Image: pw-hpc-c7-x86-64-v24-slurm

Snapshot Region: US-EAST-1

Name: wrf_cluster_demo_05

Description: Autobuild based on latest and spack and miniconda tarballs

Snapshot Build Script:

```
# Install some packages
sudo yum install -y centos-release-scl
sudo yum install -y devtoolset-7
sudo yum install -y wget git git-lfs screen zip unzip bzip2 ksh csh time psmisc gcc cmake ImageMagick gdal-python libgeotiff-
devel libtiff-devel wgrib wgrib2 python39-setuptools python39-devel python34-pip nco wgrib wgr\
ib2 ncview lapack-devel blas-devel pip awscli gcc glibc glibc-common gcc-c++ kernel-devel gc gcc++ gcc-c++ nco wgrib wgrib2
ncview bc nc jq libXScrnSaver alsa-lib xorg-x11-server-Xorg  gtk+-devel gtk2-devel

# Make  the staging ground
export STAGING_DIR=/var/lib/pworks
sudo mkdir -p $STAGING_DIR
sudo chmod a+rwx $STAGING_DIR
cd $STAGING_DIR

echo Download the tarballs...
```

Save Snapshot Config | Back

Provisioning Log:

```
941 ==> amazon-ebs.aws: Waiting
942 ==> amazon-ebs.aws: Skipping
943 ==> amazon-ebs.aws: Adding t
    023d6b1b905010ad4)...
944 ==> amazon-ebs.aws: Tagging
945 ==> amazon-ebs.aws: Creating
946     amazon-ebs.aws: Adding t
947     amazon-ebs.aws: Adding t
948     amazon-ebs.aws: Adding t
949     amazon-ebs.aws: Adding t
    0e4ab3332139fa643"
950     amazon-ebs.aws: Adding tag: "Supportfolio": ""
951 ==> amazon-ebs.aws: Creating snapshot tags
952 ==> amazon-ebs.aws: Terminating the source AWS instance...
953 ==> amazon-ebs.aws: Cleaning up any extra volumes...
954 ==> amazon-ebs.aws: No volumes to clean up, skipping
955 ==> amazon-ebs.aws: Deleting temporary security group...
956 ==> amazon-ebs.aws: Deleting temporary keypair...
957 Build 'amazon-ebs.aws' finished after 31 minutes 47 seconds.
958
959 ==> Wait completed after 31 minutes 47 seconds
960
961 ==> Builds finished. The artifacts of successful builds are:
962 --> amazon-ebs.aws: AMIs were created:
```

Top

Provision Snapshot | Delete Snapshot

First, “Create Snapshot”, then that button becomes the “Save Snapshot Config” button each time there is an update to the snapshot build script, etc.

Build scripts for WRF images are available for GCE and AWS.

Setup step 3a: Manage resources

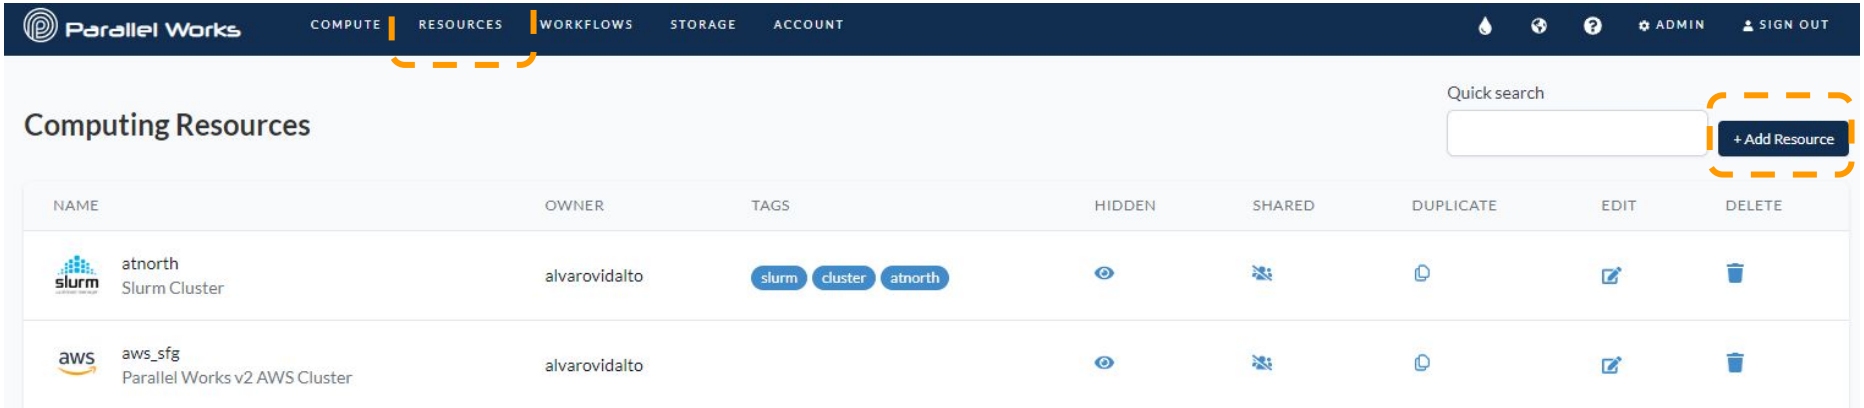

After logging in, go to the Resources Tab and select either Add Resource or an existing resource. There are three types of resources:

1. persistent clusters, e.g. atNorth (“Slurm Cluster” provider)
2. **cloud clusters (please use V2 clusters for this trial)**
3. worker pools (workers nodes are independent, no head node)

When configuring a new cluster,

1. Select a project created in step 1
2. Select an image created in step 2
3. Select the compute resources of interest. **Note:** AWS hpc6a instances are only available in us-east-2.

Setup step 3b: Configure resources

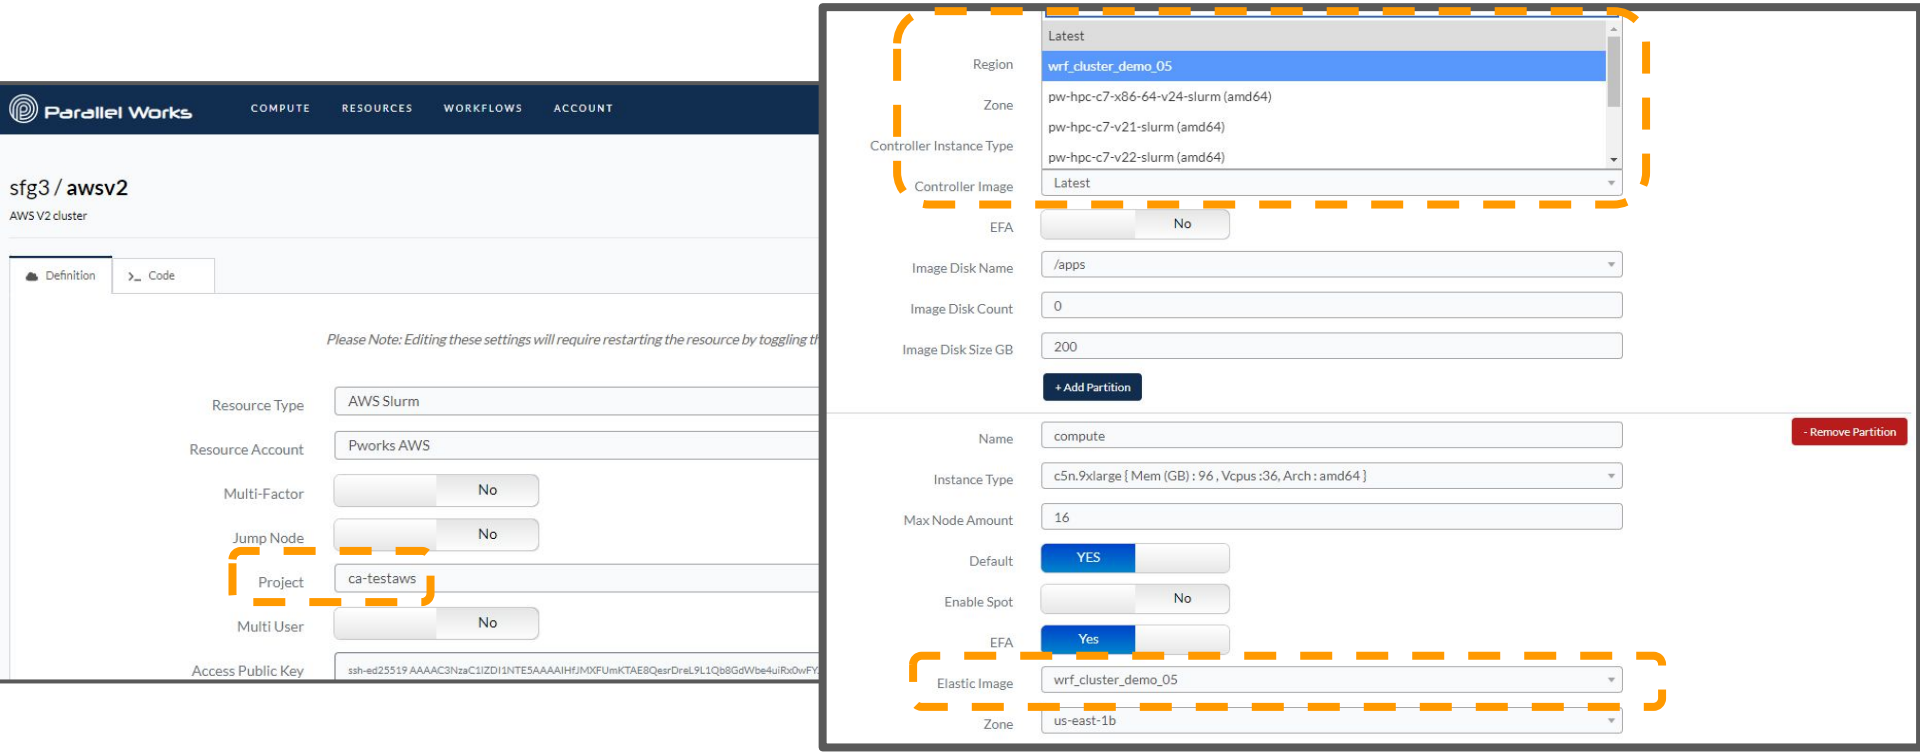

When configuring a new cluster,

1. Select a project created in step 1 (left figure)
2. Select an image created in step 2 (right figure) + select instances, etc.
3. Select the compute resources of interest. **Note:** AWS hpc6a instances are only available in us-east-2.

Example 1:

- PW workflow clones a Github repository at runtime (when a PW job is submitted)
- Github repository has two branches:
  - Main: Is cloned by default in production
  - Development: Used for development
- A Github action is used to test new releases of the development branch and merge them into the main branch
  - Github action uses PW Client to automate workflow execution across multiple resource providers
- Deploy keys are used to control read and write access to the repository
- Links to the test-workflow-action and its implementation in the weather demo repository

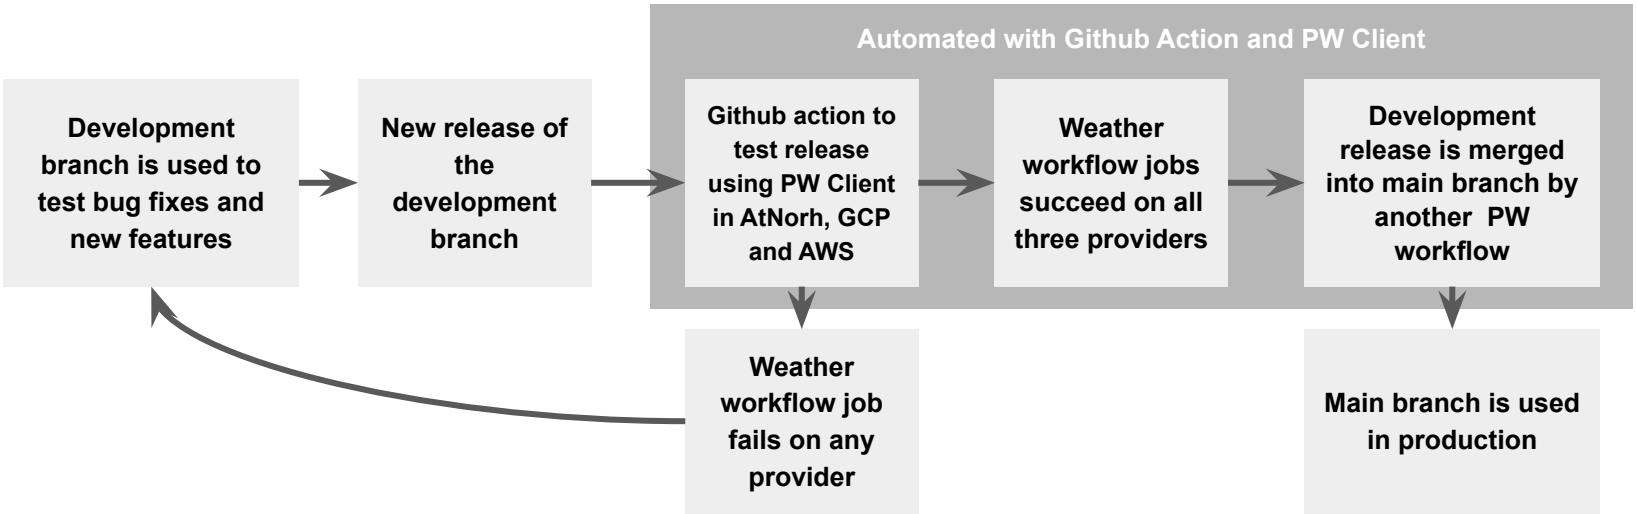

Example 1:

These are the 4 PW jobs launched by the action on the new development release:

- **56758**: Testing the weather-cluster-demo workflow in AtNorth
- **56759**: Testing the weather-cluster-demo workflow in GCP
- **56760**: Testing the weather-cluster-demo workflow in AWS
- **56761**: Merging the development release into the merge branch with the merge_github_branches workflow

Job status is “Complete” if the exit code is 0 and “Error” otherwise. Error handling (including exit code) is up to the workflow developer (see /pw/workflows/weather-cluster-demo/main.sh)

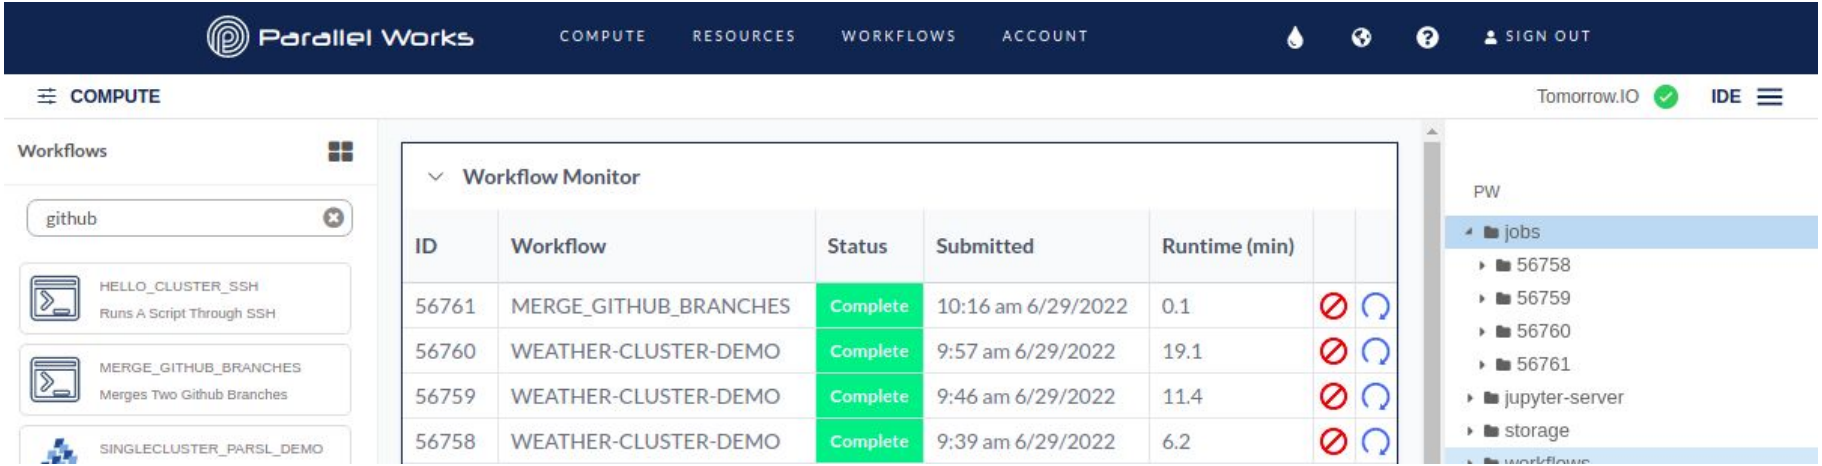

Key Components

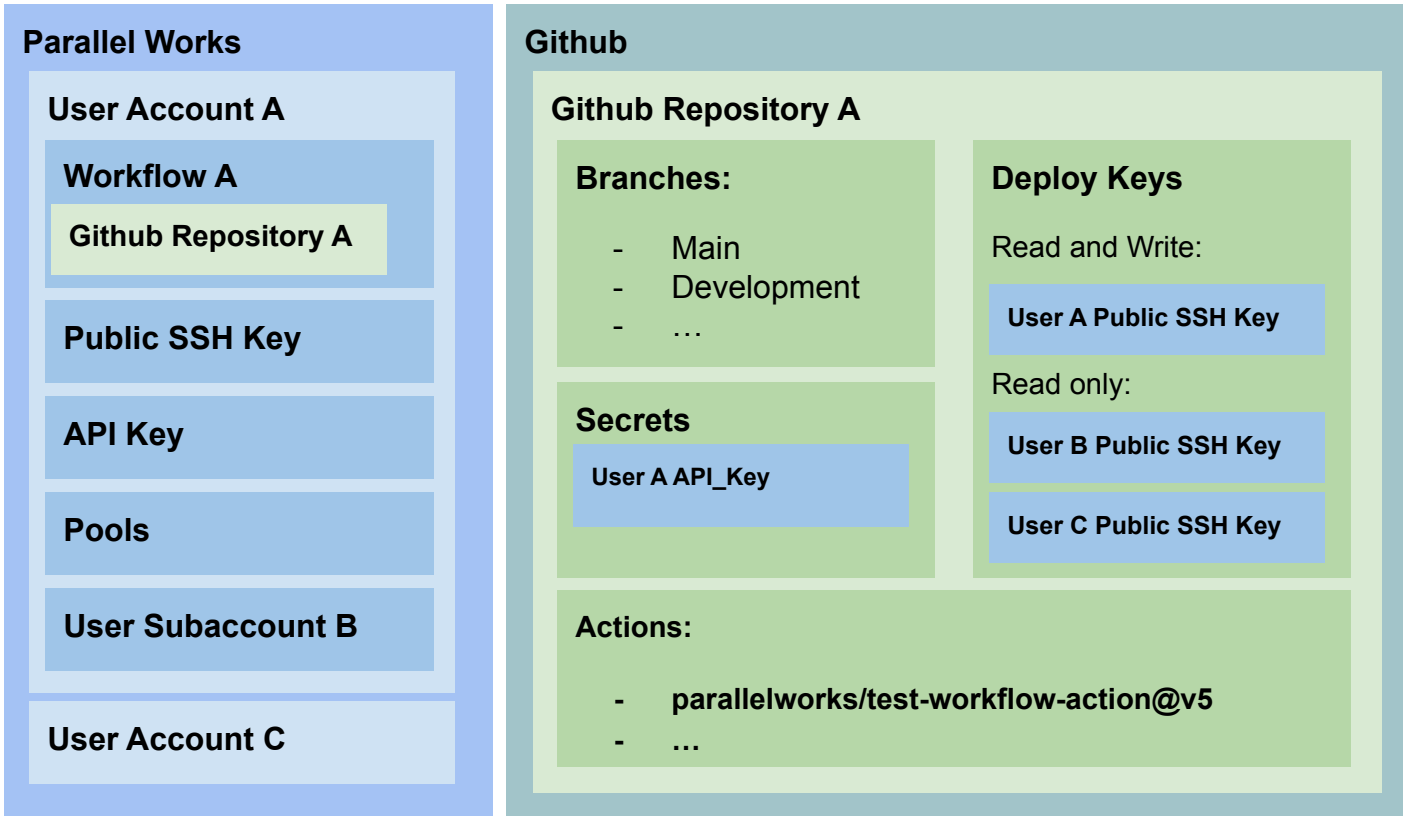

Github action runs Workflow A in User Account A

- Need User API Key

Workflow A merges development branch into main

- User A needs read and write access

Users B and C use the workflow in production

- Need read access

Sharing Workflows in PW

Can control read, write and admin access to your workflow in PW and/or use deploy keys in Github

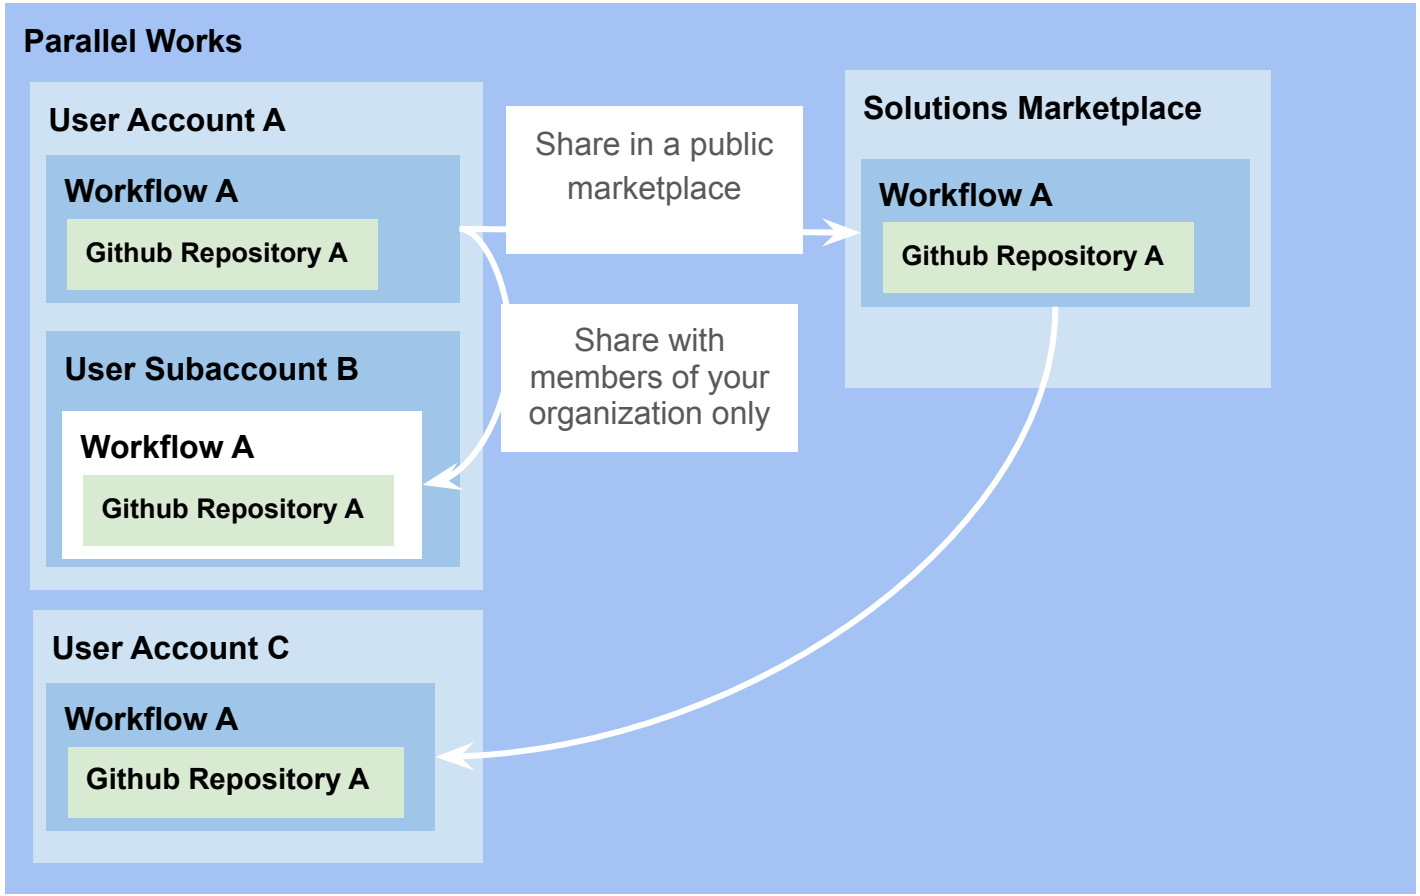

Creating Workflows in PW

These are the options to create a workflow in PW:

1. Import a workflow from the solutions marketplace
2. Duplicate an existing workflow in your account
3. Add a new workflow (not recommended)

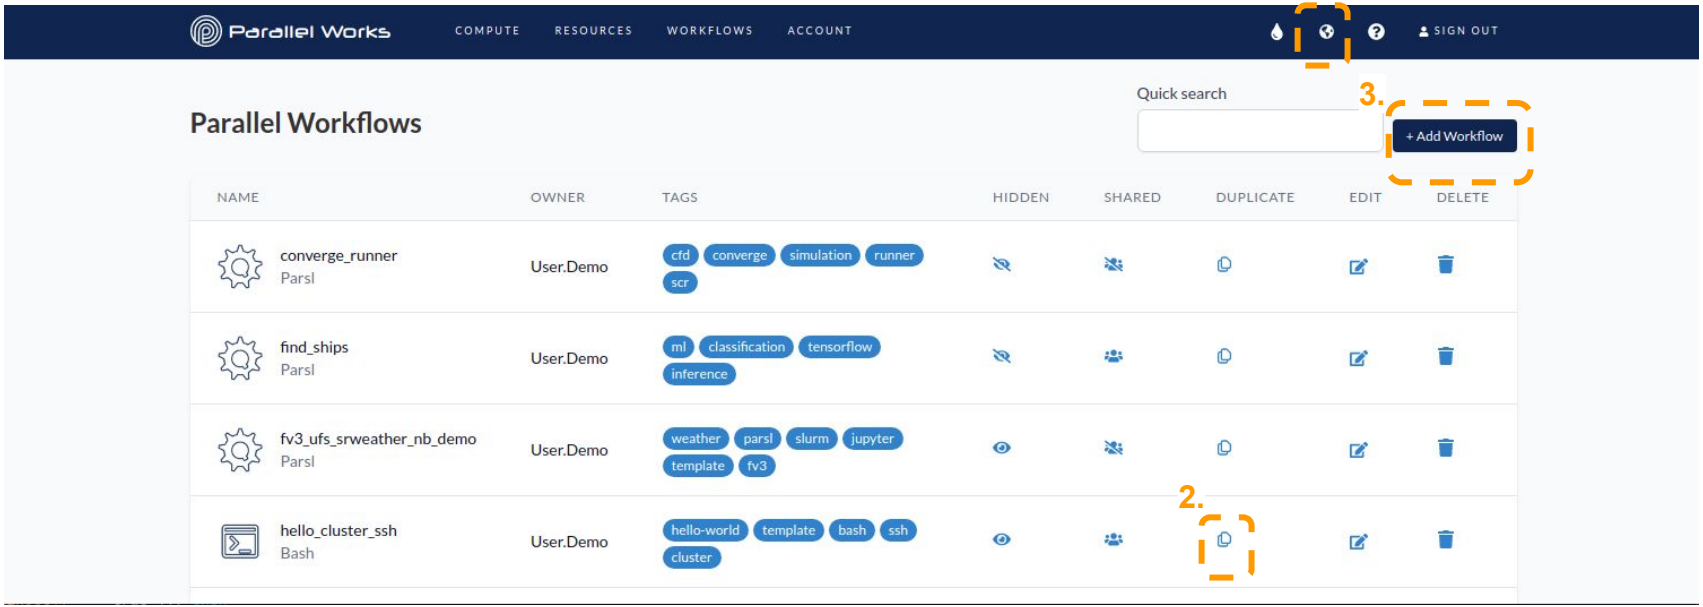

PW Jobs

When a workflow is executed a PW job is created

1. The workflow is copied to and executed in **/pw/jobs/job-number**
2. The workflow's input form, command and arguments are defined in the **/pw/workflows/workflow-name/workflow.xml** file. For example, the XML file below runs the command:

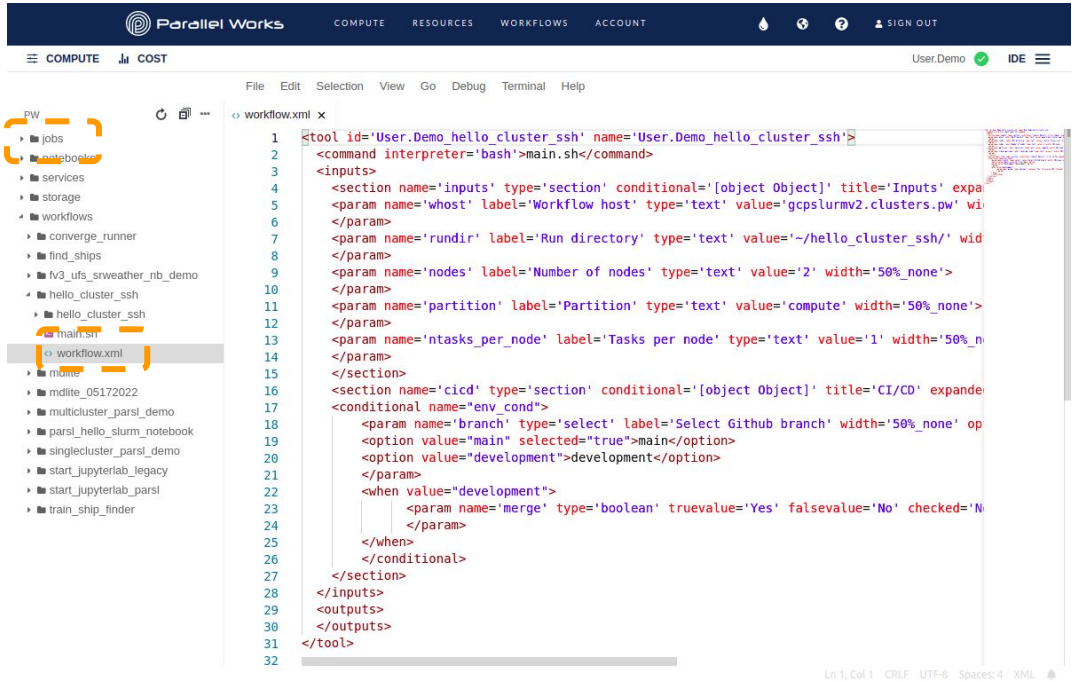


```
bash main.sh \
    --whost gcpslurmv2.clusters.pw
\
    --rundir ~/hello_cluster_ssh/ \
    --nodes 2 \
    --partition compute \
    --ntasks_per_node 1 \
    --branch main \
```

Note that these parameter values are the default values for this workflow. Users may specify their own parameter values in the input form or the PW client (see next slide)

PW Jobs

Workflows can be executed from the input form (web UI) and using the PW client (automated)

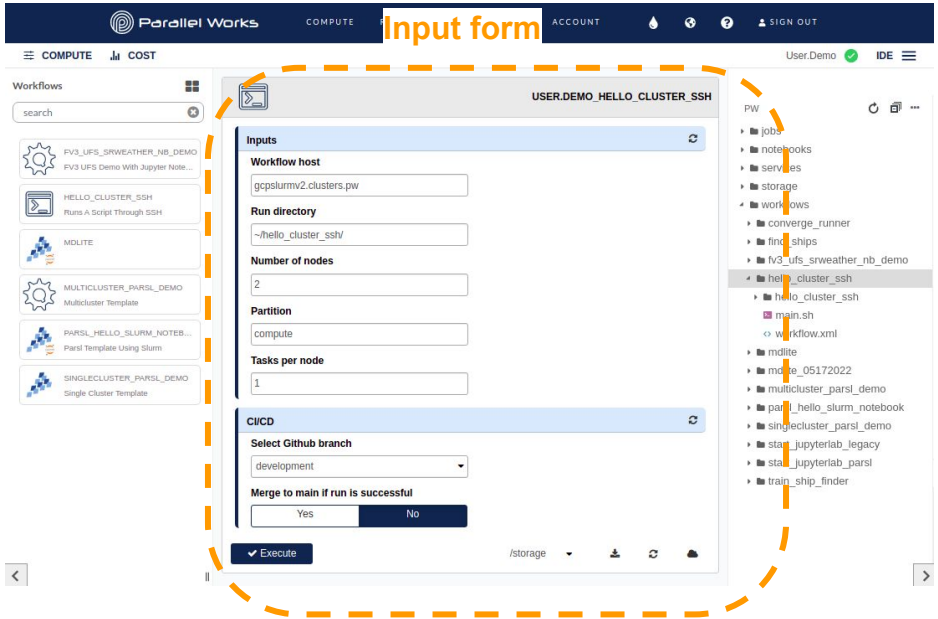

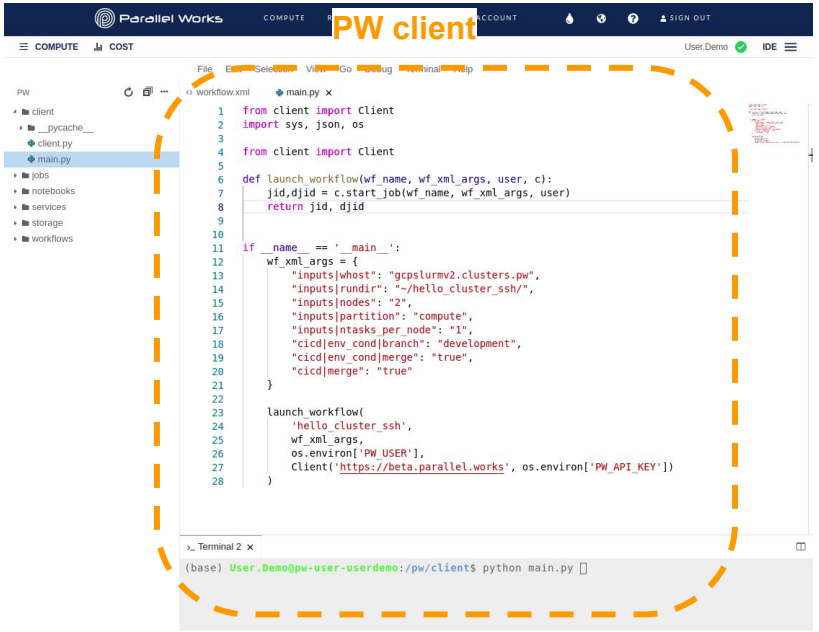

Github Deploy Keys

Use deploy keys to manage access of PW accounts to Github repositories. Follow these steps:

1. Create new ssh keys under **~/.ssh/org_name.repo_name.github.id_rsa** by running the following command in a terminal window of the PW IDE:

   ssh-keygen -t rsa

2. Create a new entry in the **~/.ssh/config** and **~/.ssh/config_custom** files:

```
Host org_name-repo_name
  HostName github.com
  User git
  IdentityFile ~/.ssh/org_name-repo_name.github.id_rsa
```

3. Add the public key to the deploy keys of the Github repository with read only or read and write permissions

4. Clone the repository with the command:

   git clone org_name-repo_name:org_name/repo_name.git

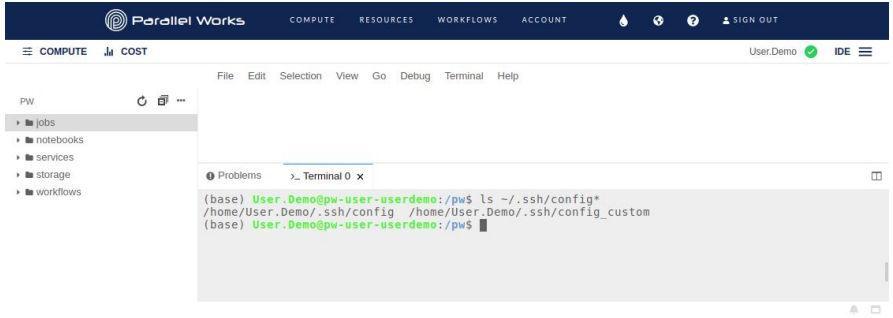

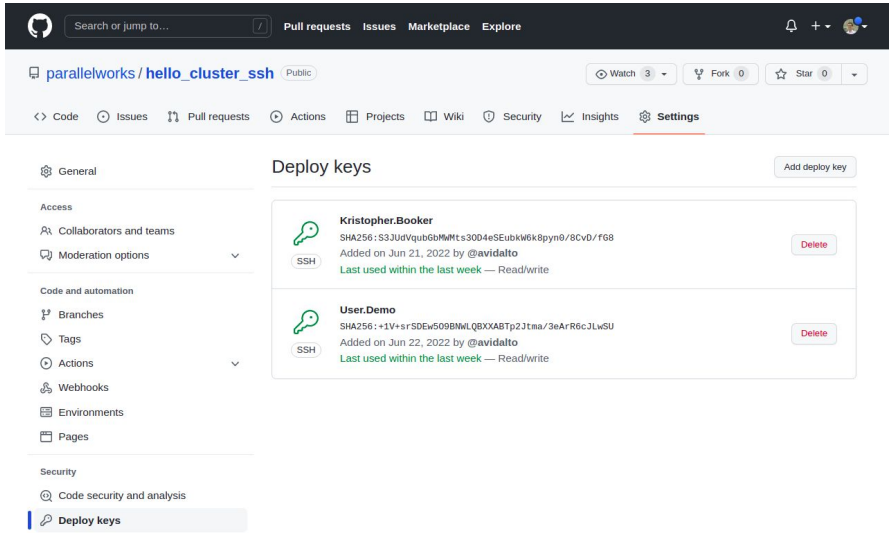

Github Actions

Use Github actions to launch PW workflows using the PW API Client. An example action is provided in the repository: https://github.com/parallelworks/test-workflow-action

As an example, this action is used in the repository:

https://github.com/parallelworks/hello_cluster_ssh

(See .github/workflows/main.yaml file)

Where the workflow-parameters are downloaded from the input form of the hello_cluster_ssh in PW

The user's API key must be added to the secrets of the repository to be used by the PW API Client. This can be found in ACCOUNT > API Key or by printing the environment variable ${PW_API_KEY}

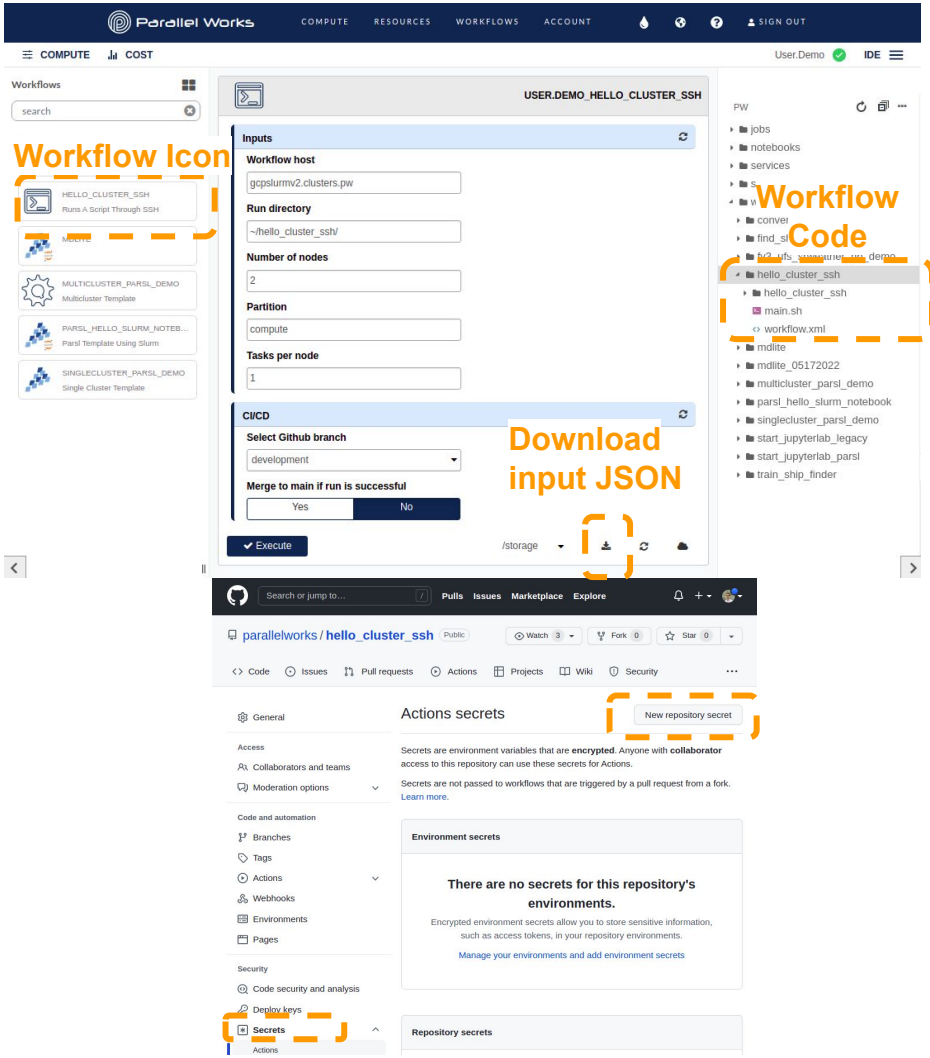

Github Calls in PW Workflows

The hello_cluster_ssh workflow is an example of:

1. Cloning a github repository every time a workflow is executed (needs read access to the repository)
2. Merging two branches (development to main) if the workflow runs successfully (needs write access to the repository)

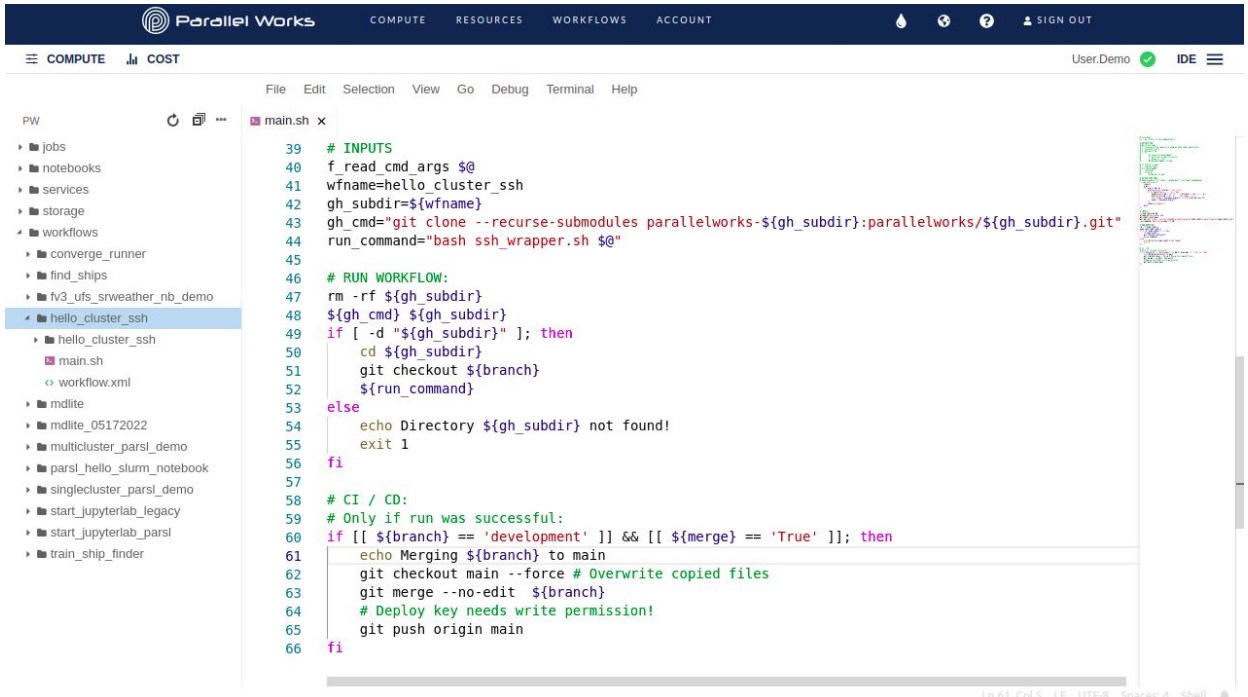

Tagging and Releasing a Repository Version

To add tag to a github repository run the following commands:

```
git tag -a -m "My new tag" vMAJOR.MINOR.PATCH
git push --follow-tags
```

Then go to Github releases (e.g.: https://github.com/parallelworks/hello_cluster_ssh/**releases**), select “draft a new release” and select your tag. This should trigger the action in the hello_cluster_ssh repository